

imph6h73h

February 9, 2024

```
[151]: # load data
from sklearn import datasets
boston = datasets.load_boston()
print(boston.keys())
```

```
dict_keys(['data', 'target', 'feature_names', 'DESCR', 'filename',
'data_module'])
```

```
/Users/prateek/anaconda3/envs/AnacondaTest/lib/python3.11/site-
packages/sklearn/utils/deprecation.py:87: FutureWarning: Function load_boston is
deprecated; `load_boston` is deprecated in 1.0 and will be removed in 1.2.
```

The Boston housing prices dataset has an ethical problem. You can refer to the documentation of this function for further details.

The scikit-learn maintainers therefore strongly discourage the use of this dataset unless the purpose of the code is to study and educate about ethical issues in data science and machine learning.

In this special case, you can fetch the dataset from the original source::

```
import pandas as pd
import numpy as np

data_url = "http://lib.stat.cmu.edu/datasets/boston"
raw_df = pd.read_csv(data_url, sep="\s+", skiprows=22, header=None)
data = np.hstack([raw_df.values[::2, :], raw_df.values[1::2, :2]])
target = raw_df.values[1::2, 2]
```

Alternative datasets include the California housing dataset (i.e. :func:`~sklearn.datasets.fetch\_california\_housing`) and the Ames housing dataset. You can load the datasets as follows::

```
from sklearn.datasets import fetch_california_housing
housing = fetch_california_housing()
```

for the California housing dataset and::

```
from sklearn.datasets import fetch_openml
housing = fetch_openml(name="house_prices", as_frame=True)
```

```
for the Ames housing dataset.
warnings.warn(msg, category=FutureWarning)
```

```
[152]: # a form of summary of the data
feature = boston.data
price = boston.target
print('data size = ', feature.shape)
print('target size = ', price.shape)
print('feature attributes: ', boston.feature_names)
print(boston.DESCR)
```

```
data size = (506, 13)
target size = (506,)
feature attributes: ['CRIM' 'ZN' 'INDUS' 'CHAS' 'NOX' 'RM' 'AGE' 'DIS' 'RAD'
'TAX' 'PTRATIO'
'B' 'LSTAT']
.. _boston_dataset:
```

Boston house prices dataset

-----

**\*\*Data Set Characteristics:\*\***

:Number of Instances: 506

:Number of Attributes: 13 numeric/categorical predictive. Median Value (attribute 14) is usually the target.

:Attribute Information (in order):

- CRIM per capita crime rate by town
- ZN proportion of residential land zoned for lots over 25,000 sq.ft.
- INDUS proportion of non-retail business acres per town
- CHAS Charles River dummy variable (= 1 if tract bounds river; 0 otherwise)
- NOX nitric oxides concentration (parts per 10 million)
- RM average number of rooms per dwelling
- AGE proportion of owner-occupied units built prior to 1940
- DIS weighted distances to five Boston employment centres
- RAD index of accessibility to radial highways
- TAX full-value property-tax rate per \$10,000
- PTRATIO pupil-teacher ratio by town
- B  $1000(B_k - 0.63)^2$  where  $B_k$  is the proportion of black people by town

- LSTAT      % lower status of the population
- MEDV      Median value of owner-occupied homes in \$1000's

:Missing Attribute Values: None

:Creator: Harrison, D. and Rubinfeld, D.L.

This is a copy of UCI ML housing dataset.

<https://archive.ics.uci.edu/ml/machine-learning-databases/housing/>

This dataset was taken from the StatLib library which is maintained at Carnegie Mellon University.

The Boston house-price data of Harrison, D. and Rubinfeld, D.L. 'Hedonic prices and the demand for clean air', J. Environ. Economics & Management, vol.5, 81-102, 1978. Used in Belsley, Kuh & Welsch, 'Regression diagnostics ...', Wiley, 1980. N.B. Various transformations are used in the table on pages 244-261 of the latter.

The Boston house-price data has been used in many machine learning papers that address regression problems.

.. topic:: References

- Belsley, Kuh & Welsch, 'Regression diagnostics: Identifying Influential Data and Sources of Collinearity', Wiley, 1980. 244-261.

- Quinlan,R. (1993). Combining Instance-Based and Model-Based Learning. In Proceedings on the Tenth International Conference of Machine Learning, 236-243, University of Massachusetts, Amherst. Morgan Kaufmann.

```
[153]: # more details of the data
import pandas as pd
df_feature = pd.DataFrame(feature, columns = boston.feature_names)
df_target = pd.DataFrame(price, columns = ['MEDV'])
df_boston = pd.concat([df_feature, df_target,], axis = 1)
```

```
[154]: df_boston.head()
```

```
[154]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	\
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	

	PTRATIO	B	LSTAT	MEDV
0	15.3	396.90	4.98	24.0
1	17.8	396.90	9.14	21.6
2	17.8	392.83	4.03	34.7
3	18.7	394.63	2.94	33.4
4	18.7	396.90	5.33	36.2

```
[155]: df_boston.describe()
```

```
[155]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM \
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000
mean	3.613524	11.363636	11.136779	0.069170	0.554695	6.284634
std	8.601545	23.322453	6.860353	0.253994	0.115878	0.702617
min	0.006320	0.000000	0.460000	0.000000	0.385000	3.561000
25%	0.082045	0.000000	5.190000	0.000000	0.449000	5.885500
50%	0.256510	0.000000	9.690000	0.000000	0.538000	6.208500
75%	3.677083	12.500000	18.100000	0.000000	0.624000	6.623500
max	88.976200	100.000000	27.740000	1.000000	0.871000	8.780000

	AGE	DIS	RAD	TAX	PTRATIO	B \
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000
mean	68.574901	3.795043	9.549407	408.237154	18.455534	356.674032
std	28.148861	2.105710	8.707259	168.537116	2.164946	91.294864
min	2.900000	1.129600	1.000000	187.000000	12.600000	0.320000
25%	45.025000	2.100175	4.000000	279.000000	17.400000	375.377500
50%	77.500000	3.207450	5.000000	330.000000	19.050000	391.440000
75%	94.075000	5.188425	24.000000	666.000000	20.200000	396.225000
max	100.000000	12.126500	24.000000	711.000000	22.000000	396.900000

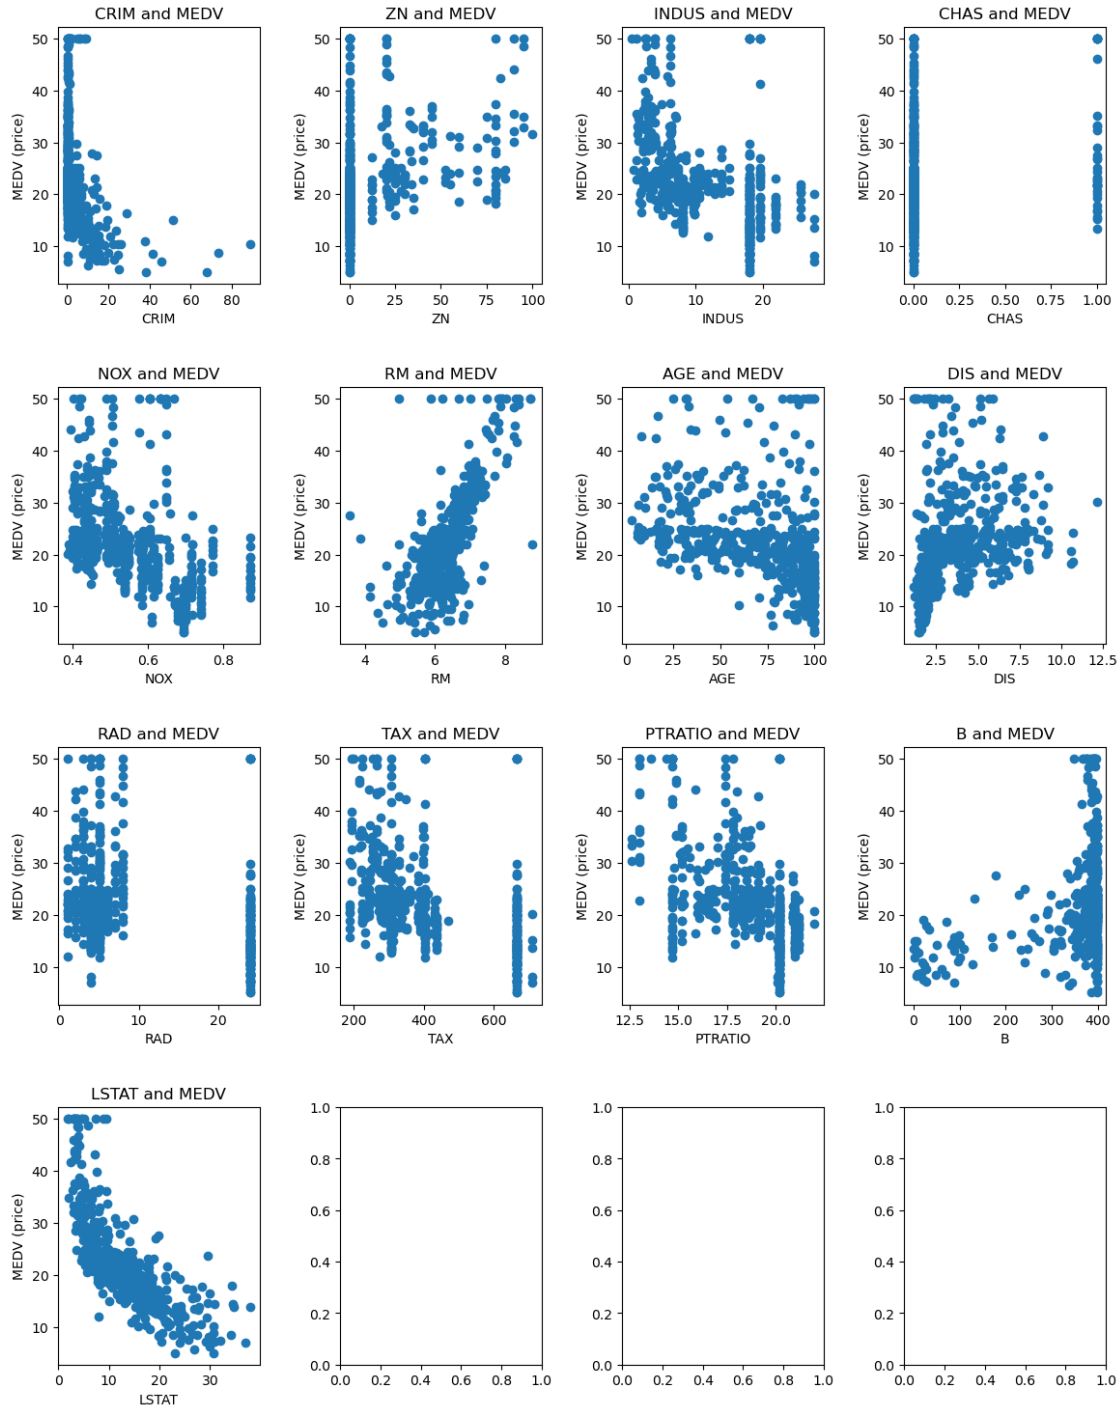
  

	LSTAT	MEDV
count	506.000000	506.000000
mean	12.653063	22.532806
std	7.141062	9.197104
min	1.730000	5.000000
25%	6.950000	17.025000
50%	11.360000	21.200000
75%	16.955000	25.000000
max	37.970000	50.000000

```
[156]: # 2.1 how does each feature relate to the price
import matplotlib.pyplot as plt
plt.figure()
fig, axes = plt.subplots(4, 4, figsize=(14, 18))
fig.subplots_adjust(wspace=.4, hspace=.4)
img_index = 0
for i in range(boston.feature_names.size):
```


```
row, col = i // 4, i % 4
axes[row][col].scatter(feature[:,i], price)
axes[row][col].set_title(boston.feature_names[i] + ' and MEDV')
axes[row][col].set_xlabel(boston.feature_names[i])
axes[row][col].set_ylabel('MEDV (price)')
plt.show()
```

<Figure size 640x480 with 0 Axes>



```
[157]: # 2.2 correlation matrix
import seaborn as sns
fig, ax = plt.subplots(figsize=(16, 10))
correlation = df_boston.corr()
sns.heatmap(correlation, annot = True, cmap = 'RdBu')
```

```
plt.show()
correlation
```



```
[157]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT	MEDV
CRIM	1.000000	-0.200469	0.406583	-0.055892	0.420972	-0.219247	0.352734	-0.379670	0.625505	0.582764	0.289946	-0.385064	0.455621	-0.388305
ZN	-0.200469	1.000000	-0.533828	-0.042697	-0.516604	0.311991	-0.569537	0.664408	-0.311948	-0.314563	-0.391679	0.175520	-0.412995	0.360445
INDUS	0.406583	-0.533828	1.000000	0.062938	0.763651	-0.391676	0.644779	-0.708027	0.595129	0.720760	0.383248	-0.356977	0.603800	-0.483725
CHAS	-0.055892	-0.042697	0.062938	1.000000	0.091203	0.091251	0.086518	-0.099176	-0.007368	-0.035587	-0.121515	0.048788	-0.053929	0.175260
NOX	0.420972	-0.516604	0.763651	0.091203	1.000000	-0.302188	0.731470	-0.769230	0.611441	0.668023	0.188933	-0.380051	0.590879	-0.427321
RM	-0.219247	0.311991	-0.391676	0.091251	-0.302188	1.000000	-0.240265	0.205246	-0.209847	-0.292048	-0.355501	0.128069	-0.613808	0.695360
AGE	0.352734	-0.569537	0.644779	0.086518	0.731470	-0.240265	1.000000	-0.747881	0.456022	0.506456	0.261515	-0.273534	0.602339	-0.376955
DIS	-0.379670	0.664408	-0.708027	-0.099176	-0.769230	0.205246	-0.747881	1.000000	-0.209847	-0.292048	-0.355501	0.128069	-0.613808	0.695360
RAD	0.625505	-0.311948	0.595129	-0.007368	0.611441	-0.209847	0.456022	-0.209847	1.000000	0.668023	0.188933	-0.380051	0.590879	-0.427321
TAX	0.582764	-0.314563	0.720760	-0.035587	0.668023	-0.292048	0.506456	-0.292048	0.668023	1.000000	-0.355501	0.128069	-0.613808	0.695360
PTRATIO	0.289946	-0.391679	0.383248	-0.121515	0.188933	-0.355501	0.261515	-0.355501	-0.355501	-0.355501	1.000000	0.128069	-0.613808	0.695360
B	-0.385064	0.175520	-0.356977	0.048788	-0.380051	0.128069	-0.273534	0.128069	-0.380051	0.128069	0.128069	1.000000	-0.613808	0.695360
LSTAT	0.455621	-0.412995	0.603800	-0.053929	0.590879	-0.613808	0.602339	-0.613808	0.590879	-0.613808	-0.613808	-0.613808	1.000000	-0.376955
MEDV	-0.388305	0.360445	-0.483725	0.175260	-0.427321	0.695360	-0.376955	0.695360	-0.427321	0.695360	0.695360	0.695360	-0.376955	1.000000

INDUS	-0.708027	0.595129	0.720760	0.383248	-0.356977	0.603800	-0.483725
CHAS	-0.099176	-0.007368	-0.035587	-0.121515	0.048788	-0.053929	0.175260
NOX	-0.769230	0.611441	0.668023	0.188933	-0.380051	0.590879	-0.427321
RM	0.205246	-0.209847	-0.292048	-0.355501	0.128069	-0.613808	0.695360
AGE	-0.747881	0.456022	0.506456	0.261515	-0.273534	0.602339	-0.376955
DIS	1.000000	-0.494588	-0.534432	-0.232471	0.291512	-0.496996	0.249929
RAD	-0.494588	1.000000	0.910228	0.464741	-0.444413	0.488676	-0.381626
TAX	-0.534432	0.910228	1.000000	0.460853	-0.441808	0.543993	-0.468536
PTRATIO	-0.232471	0.464741	0.460853	1.000000	-0.177383	0.374044	-0.507787
B	0.291512	-0.444413	-0.441808	-0.177383	1.000000	-0.366087	0.333461
LSTAT	-0.496996	0.488676	0.543993	0.374044	-0.366087	1.000000	-0.737663
MEDV	0.249929	-0.381626	-0.468536	-0.507787	0.333461	-0.737663	1.000000

```
[158]: # train-test split
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(feature, price, test_size=0.
↳3, random_state=8)
```

```
[159]: # 2.3 linear regression and ridge regression
import numpy as np

def least_square(X, y):
    #TODO
    theta = np.matmul(X.transpose(),X)
    theta = np.linalg.inv(theta)
    theta = np.matmul(theta,X.transpose())
    theta = np.matmul(theta, y)
    return theta

def ridge_reg(X, y, eta):
    #TODO
    theta = np.matmul(X.transpose(),X)
    num_rows, num_columns = theta.shape
    theta = theta + (eta/2)*np.identity(num_rows)
    theta = np.linalg.inv(theta)
    theta = np.matmul(theta,X.transpose())
    theta = np.matmul(theta, y)
    return theta

# apply linear regression
theta = least_square(X_train, y_train)
df_theta = pd.DataFrame(zip(boston.feature_names,
↳theta),columns=['Feature', 'Coeff'])
print("Linear Regression Output")
display(df_theta)

df_theta_r_temp = df_theta
```



```

# apply ridge regression
for i in range(25):
    theta_r = ridge_reg(X_train, y_train, i)
    df_theta_r = pd.DataFrame(zip(boston.feature_names,
    theta_r), columns=['Feature', 'Coeff'])
    print("Ridge Regression Output with eta = ", i)
    if(i == 20):
        df_theta_r_temp = df_theta_r
        display(df_theta_r)

# Take eta as 20, for the right balance of training vs testing dataset accuracy
df_theta_r = df_theta_r_temp

```

#### Linear Regression Output

	Feature	Coeff
0	CRIM	-0.099324
1	ZN	0.052251
2	INDUS	0.004516
3	CHAS	2.957261
4	NOX	1.127938
5	RM	5.854198
6	AGE	-0.014957
7	DIS	-0.920844
8	RAD	0.159519
9	TAX	-0.008934
10	PTRATIO	-0.435674
11	B	0.014905
12	LSTAT	-0.474751

#### Ridge Regression Output with eta = 0

	Feature	Coeff
0	CRIM	-0.099324
1	ZN	0.052251
2	INDUS	0.004516
3	CHAS	2.957261
4	NOX	1.127938
5	RM	5.854198
6	AGE	-0.014957
7	DIS	-0.920844
8	RAD	0.159519
9	TAX	-0.008934
10	PTRATIO	-0.435674
11	B	0.014905
12	LSTAT	-0.474751

#### Ridge Regression Output with eta = 1

	Feature	Coeff
--	---------	-------

0	CRIM	-0.099514
1	ZN	0.052396
2	INDUS	0.005971
3	CHAS	2.901961
4	NOX	0.928628
5	RM	5.852968
6	AGE	-0.014481
7	DIS	-0.920441
8	RAD	0.159970
9	TAX	-0.008921
10	PTRATIO	-0.433651
11	B	0.014964
12	LSTAT	-0.474664

Ridge Regression Output with eta = 2

	Feature	Coeff
0	CRIM	-0.099660
1	ZN	0.052550
2	INDUS	0.006999
3	CHAS	2.846734
4	NOX	0.810144
5	RM	5.848060
6	AGE	-0.014077
7	DIS	-0.919409
8	RAD	0.160346
9	TAX	-0.008919
10	PTRATIO	-0.431701
11	B	0.015015
12	LSTAT	-0.474969

Ridge Regression Output with eta = 3

	Feature	Coeff
0	CRIM	-0.099782
1	ZN	0.052709
2	INDUS	0.007809
3	CHAS	2.792744
4	NOX	0.731256
5	RM	5.841327
6	AGE	-0.013709
7	DIS	-0.918069
8	RAD	0.160677
9	TAX	-0.008922
10	PTRATIO	-0.429775
11	B	0.015062
12	LSTAT	-0.475464

Ridge Regression Output with eta = 4

	Feature	Coeff
--	---------	-------

0	CRIM	-0.099889
1	ZN	0.052870
2	INDUS	0.008487
3	CHAS	2.740394
4	NOX	0.674718
5	RM	5.833546
6	AGE	-0.013363
7	DIS	-0.916557
8	RAD	0.160976
9	TAX	-0.008926
10	PTRATIO	-0.427853
11	B	0.015106
12	LSTAT	-0.476063

Ridge Regression Output with eta = 5

	Feature	Coeff
0	CRIM	-0.099986
1	ZN	0.053032
2	INDUS	0.009079
3	CHAS	2.689813
4	NOX	0.632042
5	RM	5.825104
6	AGE	-0.013033
7	DIS	-0.914939
8	RAD	0.161248
9	TAX	-0.008932
10	PTRATIO	-0.425927
11	B	0.015148
12	LSTAT	-0.476724

Ridge Regression Output with eta = 6

	Feature	Coeff
0	CRIM	-0.100074
1	ZN	0.053194
2	INDUS	0.009608
3	CHAS	2.641017
4	NOX	0.598560
5	RM	5.816216
6	AGE	-0.012713
7	DIS	-0.913252
8	RAD	0.161499
9	TAX	-0.008938
10	PTRATIO	-0.423992
11	B	0.015189
12	LSTAT	-0.477424

Ridge Regression Output with eta = 7

	Feature	Coeff
--	---------	-------

0	CRIM	-0.100155
1	ZN	0.053356
2	INDUS	0.010089
3	CHAS	2.593974
4	NOX	0.571497
5	RM	5.807010
6	AGE	-0.012402
7	DIS	-0.911518
8	RAD	0.161731
9	TAX	-0.008945
10	PTRATIO	-0.422048
11	B	0.015228
12	LSTAT	-0.478151

Ridge Regression Output with eta = 8

	Feature	Coeff
0	CRIM	-0.100231
1	ZN	0.053517
2	INDUS	0.010532
3	CHAS	2.548625
4	NOX	0.549095
5	RM	5.797569
6	AGE	-0.012097
7	DIS	-0.909751
8	RAD	0.161946
9	TAX	-0.008951
10	PTRATIO	-0.420093
11	B	0.015267
12	LSTAT	-0.478894

Ridge Regression Output with eta = 9

	Feature	Coeff
0	CRIM	-0.100302
1	ZN	0.053679
2	INDUS	0.010943
3	CHAS	2.504903
4	NOX	0.530189
5	RM	5.787950
6	AGE	-0.011799
7	DIS	-0.907961
8	RAD	0.162146
9	TAX	-0.008956
10	PTRATIO	-0.418128
11	B	0.015304
12	LSTAT	-0.479650

Ridge Regression Output with eta = 10

	Feature	Coeff
--	---------	-------

0	CRIM	-0.100368
1	ZN	0.053839
2	INDUS	0.011329
3	CHAS	2.462736
4	NOX	0.513974
5	RM	5.778194
6	AGE	-0.011505
7	DIS	-0.906153
8	RAD	0.162332
9	TAX	-0.008962
10	PTRATIO	-0.416153
11	B	0.015341
12	LSTAT	-0.480413

Ridge Regression Output with eta = 11

	Feature	Coeff
0	CRIM	-0.100431
1	ZN	0.053999
2	INDUS	0.011691
3	CHAS	2.422052
4	NOX	0.499876
5	RM	5.768328
6	AGE	-0.011216
7	DIS	-0.904333
8	RAD	0.162505
9	TAX	-0.008967
10	PTRATIO	-0.414168
11	B	0.015378
12	LSTAT	-0.481181

Ridge Regression Output with eta = 12

	Feature	Coeff
0	CRIM	-0.100490
1	ZN	0.054158
2	INDUS	0.012033
3	CHAS	2.382781
4	NOX	0.487477
5	RM	5.758376
6	AGE	-0.010930
7	DIS	-0.902504
8	RAD	0.162666
9	TAX	-0.008971
10	PTRATIO	-0.412174
11	B	0.015414
12	LSTAT	-0.481953

Ridge Regression Output with eta = 13

	Feature	Coeff
--	---------	-------

0	CRIM	-0.100545
1	ZN	0.054317
2	INDUS	0.012357
3	CHAS	2.344856
4	NOX	0.476462
5	RM	5.748354
6	AGE	-0.010648
7	DIS	-0.900668
8	RAD	0.162816
9	TAX	-0.008975
10	PTRATIO	-0.410173
11	B	0.015449
12	LSTAT	-0.482726

Ridge Regression Output with eta = 14

	Feature	Coeff
0	CRIM	-0.100598
1	ZN	0.054475
2	INDUS	0.012665
3	CHAS	2.308210
4	NOX	0.466590
5	RM	5.738275
6	AGE	-0.010370
7	DIS	-0.898828
8	RAD	0.162955
9	TAX	-0.008979
10	PTRATIO	-0.408164
11	B	0.015484
12	LSTAT	-0.483500

Ridge Regression Output with eta = 15

	Feature	Coeff
0	CRIM	-0.100648
1	ZN	0.054632
2	INDUS	0.012958
3	CHAS	2.272783
4	NOX	0.457674
5	RM	5.728152
6	AGE	-0.010094
7	DIS	-0.896985
8	RAD	0.163084
9	TAX	-0.008982
10	PTRATIO	-0.406149
11	B	0.015518
12	LSTAT	-0.484274

Ridge Regression Output with eta = 16

	Feature	Coeff
--	---------	-------

0	CRIM	-0.100695
1	ZN	0.054789
2	INDUS	0.013238
3	CHAS	2.238517
4	NOX	0.449568
5	RM	5.717993
6	AGE	-0.009822
7	DIS	-0.895141
8	RAD	0.163205
9	TAX	-0.008985
10	PTRATIO	-0.404128
11	B	0.015552
12	LSTAT	-0.485046

Ridge Regression Output with eta = 17

	Feature	Coeff
0	CRIM	-0.100740
1	ZN	0.054945
2	INDUS	0.013505
3	CHAS	2.205356
4	NOX	0.442153
5	RM	5.707805
6	AGE	-0.009552
7	DIS	-0.893296
8	RAD	0.163316
9	TAX	-0.008988
10	PTRATIO	-0.402102
11	B	0.015585
12	LSTAT	-0.485818

Ridge Regression Output with eta = 18

	Feature	Coeff
0	CRIM	-0.100782
1	ZN	0.055100
2	INDUS	0.013761
3	CHAS	2.173249
4	NOX	0.435333
5	RM	5.697594
6	AGE	-0.009285
7	DIS	-0.891451
8	RAD	0.163420
9	TAX	-0.008990
10	PTRATIO	-0.400072
11	B	0.015618
12	LSTAT	-0.486588

Ridge Regression Output with eta = 19

	Feature	Coeff
--	---------	-------

0	CRIM	-0.100823
1	ZN	0.055254
2	INDUS	0.014006
3	CHAS	2.142145
4	NOX	0.429030
5	RM	5.687367
6	AGE	-0.009020
7	DIS	-0.889607
8	RAD	0.163516
9	TAX	-0.008991
10	PTRATIO	-0.398037
11	B	0.015651
12	LSTAT	-0.487356

Ridge Regression Output with eta = 20

	Feature	Coeff
0	CRIM	-0.100861
1	ZN	0.055408
2	INDUS	0.014242
3	CHAS	2.112001
4	NOX	0.423179
5	RM	5.677128
6	AGE	-0.008757
7	DIS	-0.887765
8	RAD	0.163604
9	TAX	-0.008993
10	PTRATIO	-0.396000
11	B	0.015683
12	LSTAT	-0.488121

Ridge Regression Output with eta = 21

	Feature	Coeff
0	CRIM	-0.100898
1	ZN	0.055561
2	INDUS	0.014468
3	CHAS	2.082771
4	NOX	0.417726
5	RM	5.666881
6	AGE	-0.008497
7	DIS	-0.885925
8	RAD	0.163685
9	TAX	-0.008994
10	PTRATIO	-0.393960
11	B	0.015715
12	LSTAT	-0.488884

Ridge Regression Output with eta = 22

	Feature	Coeff
--	---------	-------



0	CRIM	-0.100933
1	ZN	0.055713
2	INDUS	0.014685
3	CHAS	2.054414
4	NOX	0.412625
5	RM	5.656629
6	AGE	-0.008239
7	DIS	-0.884088
8	RAD	0.163760
9	TAX	-0.008994
10	PTRATIO	-0.391918
11	B	0.015747
12	LSTAT	-0.489645

Ridge Regression Output with eta = 23

	Feature	Coeff
0	CRIM	-0.100966
1	ZN	0.055865
2	INDUS	0.014894
3	CHAS	2.026893
4	NOX	0.407837
5	RM	5.646376
6	AGE	-0.007983
7	DIS	-0.882254
8	RAD	0.163829
9	TAX	-0.008994
10	PTRATIO	-0.389874
11	B	0.015778
12	LSTAT	-0.490402

Ridge Regression Output with eta = 24

	Feature	Coeff
0	CRIM	-0.100997
1	ZN	0.056015
2	INDUS	0.015095
3	CHAS	2.000170
4	NOX	0.403330
5	RM	5.636124
6	AGE	-0.007729
7	DIS	-0.880422
8	RAD	0.163892
9	TAX	-0.008994
10	PTRATIO	-0.387829
11	B	0.015809
12	LSTAT	-0.491157

```
[160]: # 2.4 evaluation
def pred_fn(X, theta):
    #TODO
    pred = np.matmul(X, theta)
    return pred

def root_mean_square_error(pred, y):
    #TODO
    diff_matrix = y - pred
    rmse = diff_matrix**2
    rmse = rmse.sum()
    rmse = rmse/np.size(y)
    rmse = np.sqrt(rmse)
    return rmse

pred_linear_train = pred_fn(X_train,df_theta.loc[:,"Coeff"])
pred_linear_test = pred_fn(X_test,df_theta.loc[:,"Coeff"])
rmse_linear_train = root_mean_square_error(pred_linear_train, y_train)
rmse_linear_test = root_mean_square_error(pred_linear_test, y_test)
print("Training Set RMSE of Linear Regression : ", rmse_linear_train)
print("Testing Set RMSE of Linear Regression : ", rmse_linear_test)

pred_ridge_train = pred_fn(X_train,df_theta_r.loc[:,"Coeff"])
pred_ridge_test = pred_fn(X_test,df_theta_r.loc[:,"Coeff"])
rmse_ridge_train = root_mean_square_error(pred_ridge_train, y_train)
rmse_ridge_test = root_mean_square_error(pred_ridge_test, y_test)
print("Training Set RMSE of Ridge Regression : ", rmse_ridge_train)
print("Testing Set RMSE of Ridge Regression : ", rmse_ridge_test)
```

```
Training Set RMSE of Linear Regression : 4.820626531838223
Testing Set RMSE of Linear Regression : 5.209217510530916
Training Set RMSE of Ridge Regression : 4.829777333975097
Testing Set RMSE of Ridge Regression : 5.189347305423606
```

```
[161]: # 2.5 linear models of top-3 features
X_train_top3 = np.stack((X_train[:,10],X_train[:,5],X_train[:,12])).transpose()
X_test_top3 = np.stack((X_test[:,10],X_test[:,5],X_test[:,12])).transpose()
# linear regression using top-3 features
theta_top3 = least_square(X_train_top3, y_train)
df_theta_top3 = pd.DataFrame(zip(boston.feature_names,
    ↪theta_top3),columns=['Feature','Coeff'])
pred_linear_train_top3 = pred_fn(X_train_top3,df_theta_top3.loc[:,"Coeff"])
pred_linear_test_top3 = pred_fn(X_test_top3,df_theta_top3.loc[:,"Coeff"])
rmse_linear_train_top3 = root_mean_square_error(pred_linear_train_top3, y_train)
rmse_linear_test_top3 = root_mean_square_error(pred_linear_test_top3, y_test)
print("Training Set RMSE of Linear Regression for top 3 features : ",
    ↪rmse_linear_train_top3)
```

```

print("Testing Set RMSE of Linear Regression for top 3 features : ",
      ↪rmse_linear_test_top3)
# ridge regression using top-3 features
theta_r_top3 = ridge_reg(X_train_top3, y_train, 20.0)
df_theta_r_top3 = pd.DataFrame(zip(boston.feature_names,
      ↪theta_r_top3), columns=['Feature', 'Coeff'])
pred_ridge_train_top3 = pred_fn(X_train_top3, df_theta_r_top3.loc[:, "Coeff"])
pred_ridge_test_top3 = pred_fn(X_test_top3, df_theta_r_top3.loc[:, "Coeff"])
rmse_ridge_train_top3 = root_mean_square_error(pred_ridge_train_top3, y_train)
rmse_ridge_test_top3 = root_mean_square_error(pred_ridge_test_top3, y_test)
print("Training Set RMSE of Ridge Regression for top 3 features and eta = 20.0 :
      ↪", rmse_ridge_train_top3)
print("Testing Set RMSE of Ridge Regression for top 3 features and eta = 20.0 :
      ↪", rmse_ridge_test_top3)

```

```

Training Set RMSE of Linear Regression for top 3 features : 5.273361751695365
Testing Set RMSE of Linear Regression for top 3 features : 5.494723646664577
Training Set RMSE of Ridge Regression for top 3 features and eta = 20.0 :
5.276310228536866
Testing Set RMSE of Ridge Regression for top 3 features and eta = 20.0 :
5.477573443118745

```

[ ]: